

ROHAN NYATI

500075940

R177219148

BATCH-5 (Ai&ML)

Experiment – 1

What is Regression?

Regression analysis is a statistical method to model the relationship between a dependent (target) and independent (predictor) variables with one or more independent variables. More specifically, Regression analysis helps us to understand how the value of the dependent variable is changing corresponding to an independent variable when other independent variables are held fixed. It predicts continuous/real values such as **temperature, age, salary, price**, etc.

Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%202%20-%20Regression/Section%204%20-%20Simple_Linear_Regression simple_linear_regression - Jupyter x

localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%202%20-%20Regression/Section%204%20-%20Simple_Linear_Regression simple_linear_regression Last Checkpoint: a few seconds ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Simple Linear Regression

Importing the libraries

```
In [0]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Importing the dataset

```
In [0]: dataset = pd.read_csv('Salary_Data.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

Splitting the dataset into the Training set and Test set

Type here to search

Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%202%20-%20Regression/Section%204%20-%20Simple_Linear_Regression simple_linear_regression - Jupyter x

localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%202%20-%20Regression/Section%204%20-%20Simple_Linear_Regression simple_linear_regression Last Checkpoint: a minute ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Importing the dataset

```
In [0]: dataset = pd.read_csv('Salary_Data.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

Splitting the dataset into the Training set and Test set

```
In [0]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/3, random_state = 0)
```

Training the Simple Linear Regression model on the Training set

```
In [4]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

Type here to search

Jupyter simple_linear_regression Last Checkpoint: a minute ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

```
In [4]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)

Out[4]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

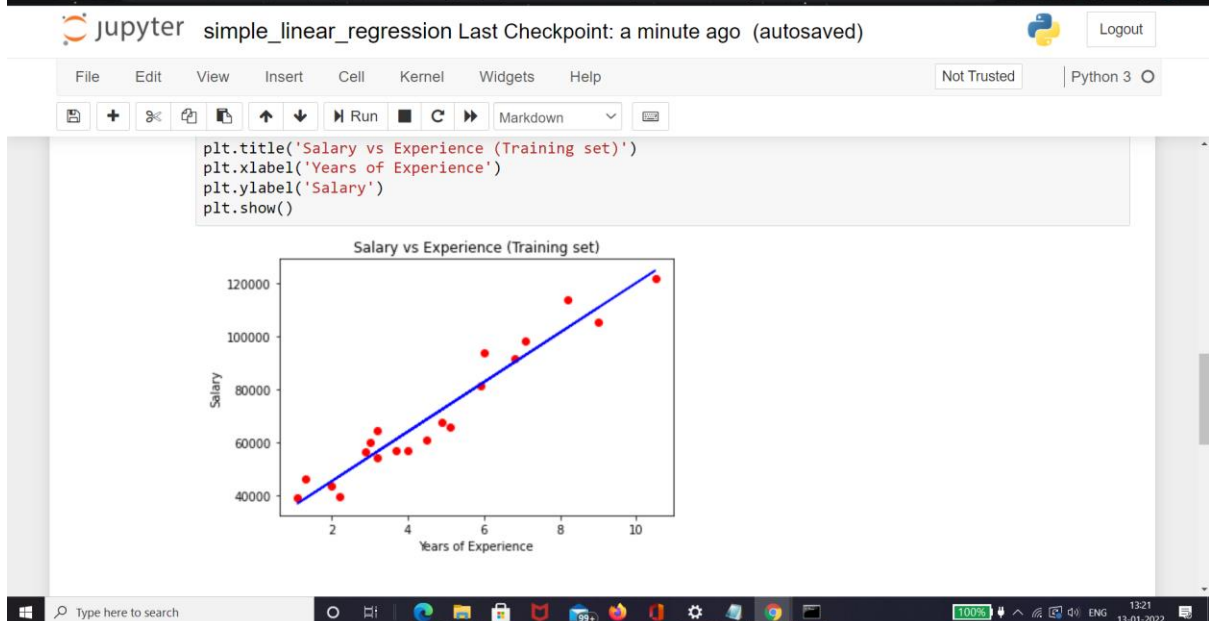
Predicting the Test set results

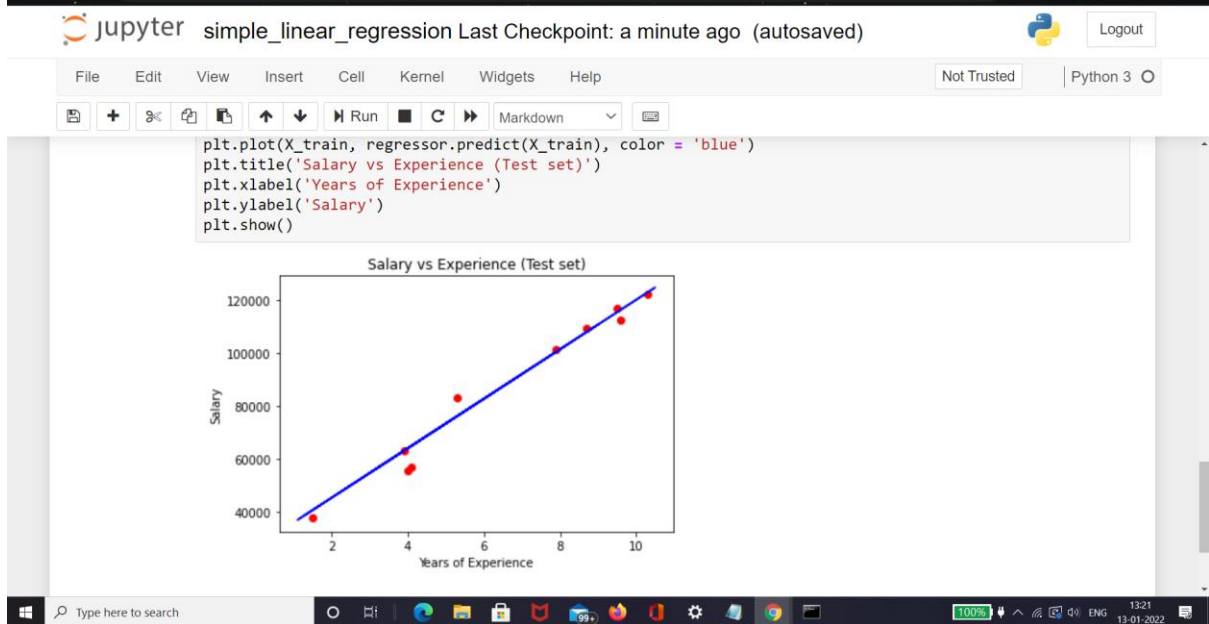
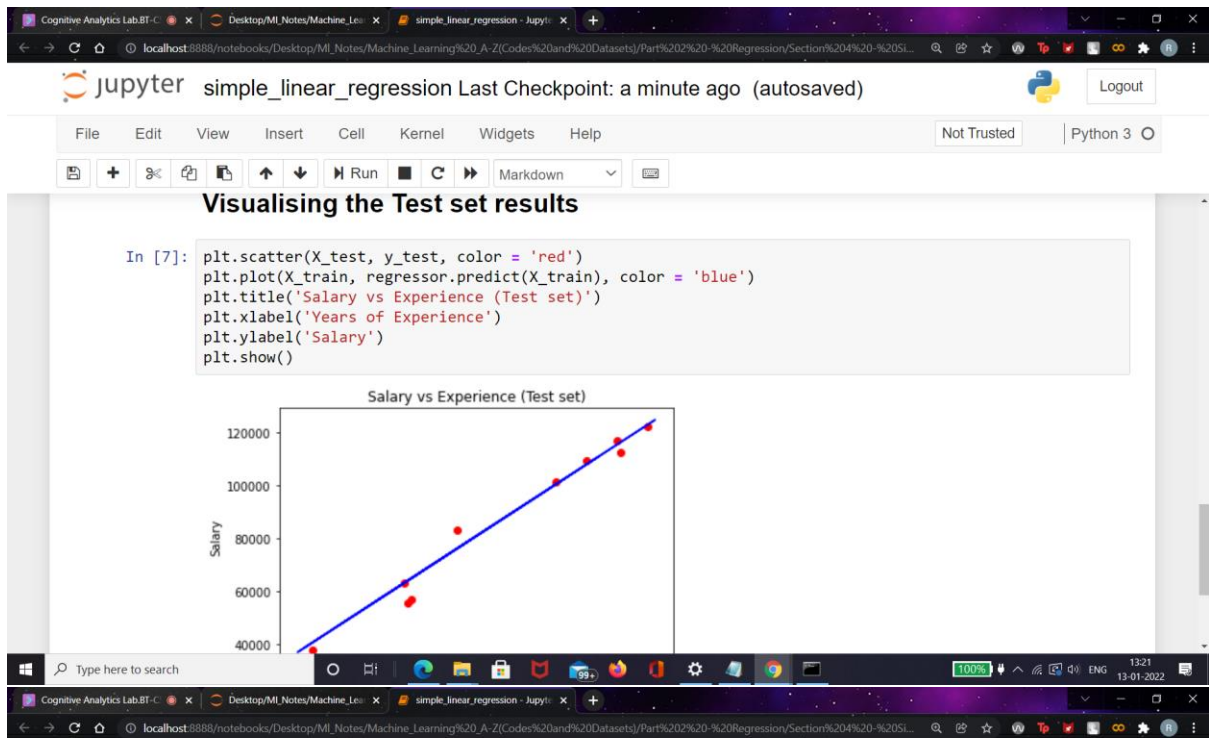
```
In [0]: y_pred = regressor.predict(X_test)
```

Visualising the Training set results

```
In [6]: plt.scatter(X_train, y_train, color = 'red')
plt.plot(X_train, regressor.predict(X_train), color = 'blue')
plt.title('Salary vs Experience (Training set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```

Windows search bar: Type here to search. Taskbar shows various applications and system tray icons including network, volume, and date/time (13:21 13-01-2022).





Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20Logistic Regression - Jupyter Notebook x simple_linear_regression - Jupyter Notebook x +

localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20Logistic Regression - Jupyter Notebook

jupyter logistic_regression (unsaved changes) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Logistic Regression

Importing the libraries

```
In [0]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Importing the dataset

```
In [0]: dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

Splitting the dataset into the Training set and Test set

Type here to search

100% 13:22 13-01-2022

jupyter logistic_regression (unsaved changes) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Splitting the dataset into the Training set and Test set

```
In [0]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
```

```
In [0]: print(X_train)
```

```
[[ 44 39000]
 [ 32 120000]
 [ 38 50000]
 [ 32 135000]
 [ 52 21000]
 [ 53 104000]
 [ 39 42000]
 [ 38 61000]
 [ 36 50000]
 [ 36 63000]
 [ 35 25000]
 [ 35 50000]
 [ 42 73000]
 [ 47 49000]
 [ 59 20000]]
```

Type here to search

100% 13:22 13-01-2022

```
Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20Logistic_Regression/Logistic_Regression.ipynb x simple_linear_regression - Jupyter Notebook x +
localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20Logistic_Regression/Logistic_Regression.ipynb
jupyter logistic_regression (unsaved changes) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3
Run
[ 46 82000]
In [0]: print(y_train)
[0 1 0 1 1 1 1 0 0 0 0 0 0 1 1 1 0 1 0 0 1 0 1 0 1 0 0 1 1 1 1 0 1 0 1 0 0 1
0 0 1 0 0 0 0 0 0 1 1 1 1 0 0 0 1 0 1 0 1 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0 1 0 1
1 1 0 0 1 1 0 0 1 1 0 1 0 0 1 1 0 1 1 1 0 0 0 0 0 1 0 0 1 1 1 1 1 0 1 1 0
1 0 0 0 0 0 0 0 1 1 0 0 1 0 0 1 0 0 0 1 0 1 1 0 1 0 0 0 0 1 0 0 0 1 1 0 0
0 0 1 0 1 0 0 0 1 0 0 0 0 1 1 1 0 0 0 0 0 0 1 1 1 1 0 1 0 0 0 0 0 1 0 0
0 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 1 1 1 0 0 0 0
0 1 1 0 0 0 0 1 0 0 0 0 1 0 1 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 1 1 0 0 0
0 0 1 0 1 1 0 0 0 0 0 1 0 1 0 0 1 0 0 1 0 1 0 0 0 0 0 1 1 1 1 0 0 0 0 1
0 0 0 0]

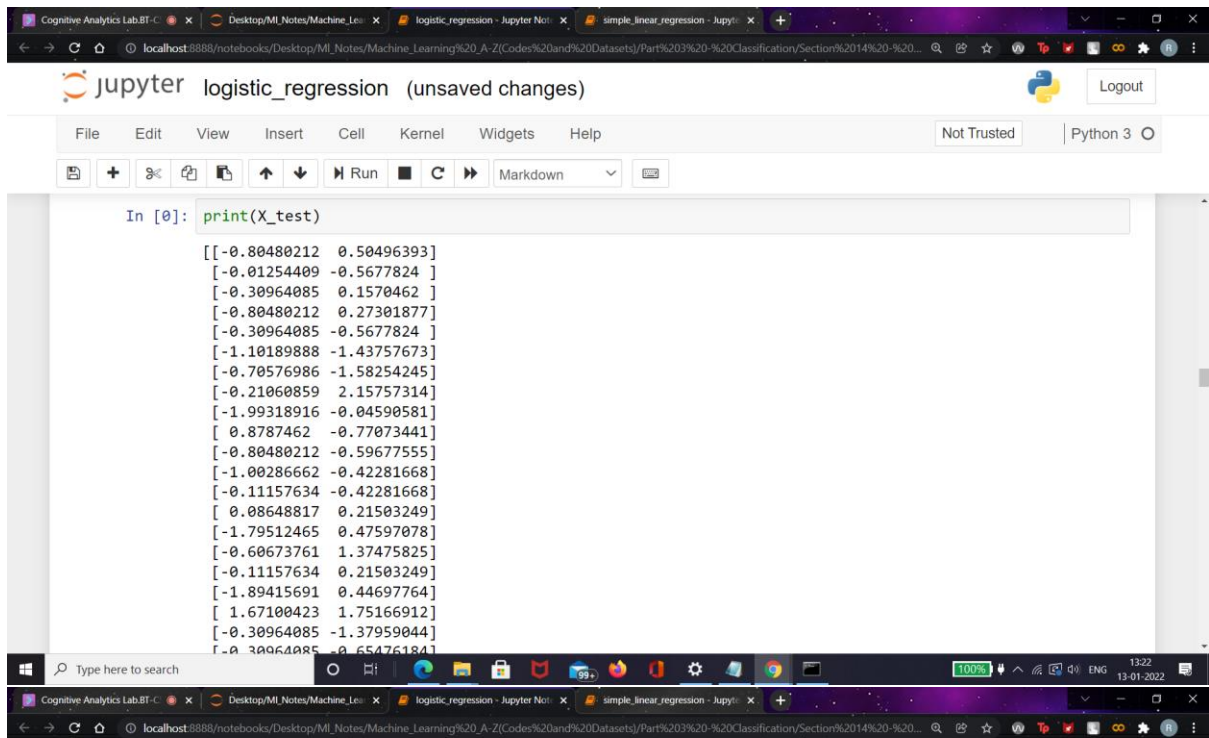
In [0]: print(X_test)
[[ 30 87000]
[ 38 50000]
[ 35 75000]
[ 30 79000]
[ 35 50000]
[ 27 20000]
```

```
jupyter logistic_regression (unsaved changes) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3
Run
[ 42 104000]]
In [0]: print(y_test)
[0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 1 0 0 0 0
0 0 1 0 0 0 0 1 0 0 1 0 1 1 0 0 0 1 1 0 0 1 0 0 1 0 0 1 0 1 0 1 0 0 0 0 1 0 0 1
0 0 0 0 1 1 1 0 0 0 1 1 0 1 1 0 0 1 0 0 0 1 0 1 1 1]

Feature Scaling

In [0]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

In [0]: print(X_train)
[[ 0.58164944 -0.88670699]
[-0.60673761  1.46173768]
[-0.01254409 -0.5677824 ]
[-0.60673761  1.89663484]
```

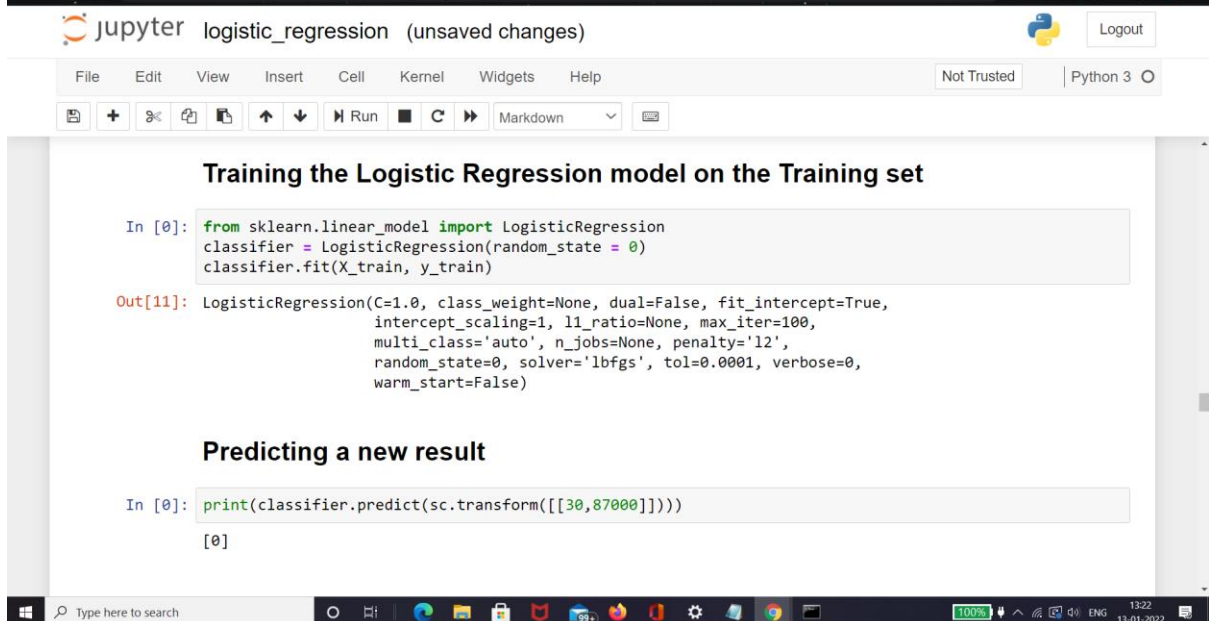



The screenshot shows a Jupyter Notebook interface with a browser window at the top. The notebook title is "logistic_regression (unsaved changes)". The menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. The status bar shows "Not Trusted" and "Python 3". The code cell contains the following Python code:

```
In [0]: print(X_test)
```

The output of the code is a 20x2 array of floating-point numbers:

```
[[-0.80480212  0.50496393]
 [-0.01254409 -0.5677824 ]
 [-0.30964085  0.1570462 ]
 [-0.80480212  0.27301877]
 [-0.30964085 -0.5677824 ]
 [-1.10189888 -1.43757673]
 [-0.70576986 -1.58254245]
 [-0.21060859  2.15757314]
 [-1.99318916 -0.04590581]
 [ 0.8787462  -0.77073441]
 [-0.80480212 -0.59677555]
 [-1.00286662 -0.42281668]
 [-0.11157634 -0.42281668]
 [ 0.08648817  0.21503249]
 [-1.79512465  0.47597078]
 [-0.60673761  1.37475825]
 [-0.11157634  0.21503249]
 [-1.89415691  0.44697764]
 [ 1.67100423  1.75166912]
 [-0.30964085 -1.37959044]
 [-0.30964085 -1.37959044]]
```



The screenshot shows a Jupyter Notebook interface with a browser window at the top. The notebook title is "logistic_regression (unsaved changes)". The menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. The status bar shows "Not Trusted" and "Python 3". The code cell contains the following Python code:

```
In [0]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)
```

The output of the code is:

```
Out[11]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=0, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

The code cell contains the following Python code:

```
In [0]: print(classifier.predict(sc.transform([[30,87000]])))
```

The output of the code is:

```
[0]
```

Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20... x

localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20...

jupyter logistic_regression (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Predicting a new result

```
In [0]: print(classifier.predict(sc.transform([[30,87000]])))
```

[0]

Predicting the Test set results

```
In [0]: y_pred = classifier.predict(X_test)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1))
```

```
[[0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [1 1]
 [0 0]]
```

Type here to search

100% 13:22 13-01-2022

Cognitive Analytics Lab.BT-C x Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20... x

localhost:8888/notebooks/Desktop/ML_Notes/Machine_Learning%20A-Z/Codes%20and%20Datasets/Part%203%20-%20Classification/Section%2014%20-%20...

jupyter logistic_regression (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Making the Confusion Matrix

```
In [0]: from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[65  3]
 [ 8 24]]
```

Out[14]: 0.89

Visualising the Training set results

```
In [0]: from matplotlib.colors import ListedColormap
X_set, y_set = sc.inverse_transform(X_train), y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 10, stop = X_set[:, 0].max() + 10, step = 0.5),
                     np.arange(start = X_set[:, 1].min() - 1000, stop = X_set[:, 1].max() + 1000, step = 1000))
```

Type here to search

100% 13:22 13-01-2022

Jupyter logistic_regression (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Out[14]: 0.89

Visualising the Training set results


```
In [0]: from matplotlib.colors import ListedColormap
X_set, y_set = sc.inverse_transform(X_train), y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 10, stop = X_set[:, 0].max() + 10, step = 0.5),
                     np.arange(start = X_set[:, 1].min() - 1000, stop = X_set[:, 1].max() + 1000, step = 1000))
plt.contourf(X1, X2, classifier.predict(sc.transform(np.array([X1.ravel(), X2.ravel()]).T)).reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1], c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Logistic Regression (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping

Jupyter logistic_regression (unsaved changes)



File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Logistic Regression (Training set)





Visualising the Test set results

```
In [0]: from matplotlib.colors import ListedColormap
```



Visualising the Test set results

```
In [0]: from matplotlib.colors import ListedColormap
X_set, y_set = sc.inverse_transform(X_test), y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 10, stop = X_set[:, 0].max() + 10, step = 0.5),
                     np.arange(start = X_set[:, 1].min() - 1000, stop = X_set[:, 1].max() + 1000, step = 1000))
plt.contourf(X1, X2, classifier.predict(sc.transform(np.array([X1.ravel(), X2.ravel()])).T).reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1], c = ListedColormap(('red', 'green'))(i), label=j)
plt.title('Logistic Regression (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```



gle row if you really want to specify the same RGB or RGBA value for all points.
'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

