

Rohan Nyati

R177219148

500075940

Btech CSE AIML B5

Neural Networks Lab 6

Installation

To install tensorflow in your linux system use the following commands.

```
sudo apt update
```

```
sudo apt install python3-dev python3-pip python3-venv
```

```
pip3 install --user --upgrade tensorflow
```

Tensors

TensorFlow operates on multidimensional arrays or *tensors* represented as `tf.Tensor` objects. Here is a two-dimensional tensor:

```
In [1]: import tensorflow as tf

x = tf.constant([[1., 2., 3.],
                 [4., 5., 6.]])

print(x)
print(x.shape)
print(x.dtype)

tf.Tensor(
[[1. 2. 3.]
 [4. 5. 6.]], shape=(2, 3), dtype=float32)
(2, 3)
<dtype: 'float32'>
```

The most important attributes of a `tf.Tensor` are its `shape` and `dtype` :

- `Tensor.shape` : tells you the size of the tensor along each of its axes.
- `Tensor.dtype` : tells you the type of all the elements in the tensor.

TensorFlow implements standard mathematical operations on tensors, as well as many operations specialized for machine learning.

For example:

```
In [2]: x + x
```

```
Out[2]: <tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[ 2.,  4.,  6.],
       [ 8., 10., 12.]], dtype=float32)>
```

```
In [3]: 5 * x
```

```
Out[3]: <tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[ 5., 10., 15.],
       [20., 25., 30.]], dtype=float32)>
```

```
In [4]: x @ tf.transpose(x)
```

```
Out[4]: <tf.Tensor: shape=(2, 2), dtype=float32, numpy=
array([[14., 32.],
       [32., 77.]], dtype=float32)>
```

```
In [5]: tf.concat([x, x, x], axis=0)
```

```
Out[5]: <tf.Tensor: shape=(6, 3), dtype=float32, numpy=
array([[1., 2., 3.],
       [4., 5., 6.],
       [1., 2., 3.],
       [4., 5., 6.],
       [1., 2., 3.],
       [4., 5., 6.]], dtype=float32)>
```

```
In [6]: tf.nn.softmax(x, axis=-1)
```

```
Out[6]: <tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[0.09003057, 0.24472848, 0.66524094],
       [0.09003057, 0.24472848, 0.66524094]], dtype=float32)>
```

```
In [7]: tf.reduce_sum(x)
```

```
Out[7]: <tf.Tensor: shape=(), dtype=float32, numpy=21.0>
```

Running large calculations on CPU can be slow. When properly configured, TensorFlow can use accelerator hardware like GPUs to execute operations very quickly.

```
In [8]: if tf.config.list_physical_devices('GPU'):
        print("TensorFlow **IS** using the GPU")
        else:
        print("TensorFlow **IS NOT** using the GPU")
```

```
TensorFlow **IS NOT** using the GPU
```

Variables

Normal `tf.Tensor` objects are immutable. To store model weights (or other mutable state) in TensorFlow use a `tf.Variable`.

```
In [9]: var = tf.Variable([0.0, 0.0, 0.0])
```

```
In [10]: var.assign([1, 2, 3])
```

```
Out[10]: <tf.Variable 'UnreadVariable' shape=(3,) dtype=float32, numpy=array([1., 2., 3.], dtype=float32)>
```

```
In [11]: var.assign_add([1, 1, 1])
```

```
Out[11]: <tf.Variable 'UnreadVariable' shape=(3,) dtype=float32, numpy=array([2., 3., 4.], dtype=float32)>
```

Automatic differentiation

Gradient descent and related algorithms are a cornerstone of modern machine learning.

To enable this, TensorFlow implements automatic differentiation (autodiff), which uses calculus to compute gradients. Typically you'll use this to calculate the gradient of a model's *error* or *loss* with respect to its weights.

```
In [12]: x = tf.Variable(1.0)
```

```
def f(x):  
    y = x**2 + 2*x - 5  
    return y
```

```
In [13]: f(x)
```

```
Out[13]: <tf.Tensor: shape=(), dtype=float32, numpy=-2.0>
```

At $x = 1.0$, $y = f(x) = (1^2 + 2 \cdot 1 - 5) = -2$.

The derivative of y is $y' = f'(x) = (2x + 2) = 4$. TensorFlow can calculate this automatically:

```
In [14]: with tf.GradientTape() as tape:
          y = f(x)

          g_x = tape.gradient(y, x) # g(x) = dy/dx

          g_x
```

```
Out[14]: <tf.Tensor: shape=(), dtype=float32, numpy=4.0>
```

This simplified example only takes the derivative with respect to a single scalar (x), but TensorFlow can compute the gradient with respect to any number of non-scalar tensors simultaneously.

Graphs and `tf.function`

While you can use TensorFlow interactively like any Python library, TensorFlow also provides tools for:

- **Performance optimization:** to speed up training and inference.
- **Export:** so you can save your model when it's done training.

These require that you use `tf.function` to separate your pure-TensorFlow code from Python.

```
In [15]: @tf.function
          def my_func(x):
              print('Tracing.\n')
              return tf.reduce_sum(x)
```

The first time you run the `tf.function`, although it executes in Python, it captures a complete, optimized graph representing the TensorFlow computations done within the function.

```
In [16]: x = tf.constant([1, 2, 3])  
my_func(x)
```

Tracing.

```
Out[16]: <tf.Tensor: shape=(), dtype=int32, numpy=6>
```

On subsequent calls TensorFlow only executes the optimized graph, skipping any non-TensorFlow steps. Below, note that `my_func` doesn't print *tracing* since `print` is a Python function, not a TensorFlow function.

```
In [17]: x = tf.constant([10, 9, 8])  
my_func(x)
```

```
Out[17]: <tf.Tensor: shape=(), dtype=int32, numpy=27>
```

A graph may not be reusable for inputs with a different *signature* (`shape` and `dtype`), so a new graph is generated instead:

```
In [18]: x = tf.constant([10.0, 9.1, 8.2], dtype=tf.float32)  
my_func(x)
```

Tracing.

```
Out[18]: <tf.Tensor: shape=(), dtype=float32, numpy=27.3>
```

Modules, layers, and models

`tf.Module` is a class for managing your `tf.Variable` objects, and the `tf.function` objects that operate on them. The `tf.Module` class is necessary to support two significant features:

1. You can save and restore the values of your variables using `tf.train.Checkpoint`. This is useful during training as it is quick to save and restore a model's state.
2. You can import and export the `tf.Variable` values *and* the `tf.function` graphs using `tf.saved_model`. This allows you to run your model independently of the Python program that created it.

Here is a complete example exporting a simple `tf.Module` object:

```
In [19]: class MyModule(tf.Module):
         def __init__(self, value):
             self.weight = tf.Variable(value)

         @tf.function
         def multiply(self, x):
             return x * self.weight
```

```
In [20]: mod = MyModule(3)
         mod.multiply(tf.constant([1, 2, 3]))
```

```
Out[20]: <tf.Tensor: shape=(3,), dtype=int32, numpy=array([3, 6, 9], dtype=int32)>
```

Save the Module :

```
In [21]: save_path = './saved'
         tf.saved_model.save(mod, save_path)
```

```
INFO:tensorflow:Assets written to: ./saved/assets
```

```
In [22]: reloaded = tf.saved_model.load(save_path)
         reloaded.multiply(tf.constant([1, 2, 3]))
```

```
Out[22]: <tf.Tensor: shape=(3,), dtype=int32, numpy=array([3, 6, 9], dtype=int32)>
```

The `tf.keras.layers.Layer` and `tf.keras.Model` classes build on `tf.Module` providing additional functionality and convenience methods for building, training, and saving models. Some of these are demonstrated in the next section.

Training loops

Now put this all together to build a basic model and train it from scratch.

First, create some example data. This generates a cloud of points that loosely follows a quadratic curve:

```
In [23]: import matplotlib
         from matplotlib import pyplot as plt

         matplotlib.rcParams['figure.figsize'] = [9, 6]
```

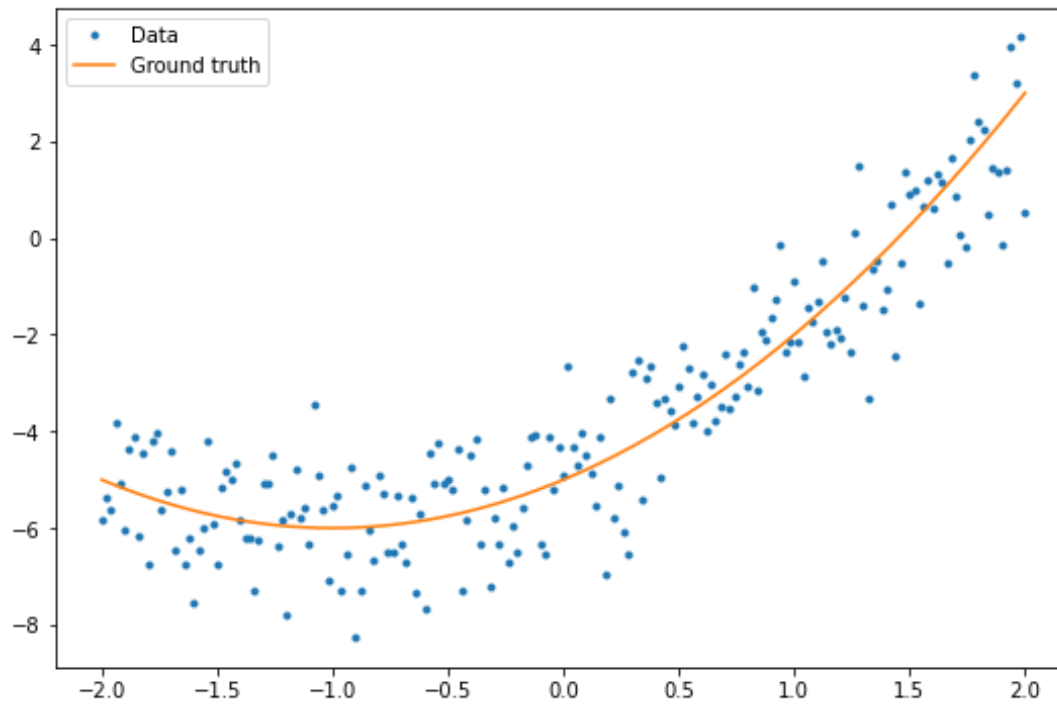


```
In [24]: x = tf.linspace(-2, 2, 201)
x = tf.cast(x, tf.float32)

def f(x):
    y = x**2 + 2*x - 5
    return y

y = f(x) + tf.random.normal(shape=[201])

plt.plot(x.numpy(), y.numpy(), '.', label='Data')
plt.plot(x, f(x), label='Ground truth')
plt.legend();
```




Create a model:

```
In [25]: class Model(tf.keras.Model):  
    def __init__(self, units):  
        super().__init__()  
        self.dense1 = tf.keras.layers.Dense(units=units,  
                                              activation=tf.nn.relu,  
                                              kernel_initializer=tf.random.normal,  
                                              bias_initializer=tf.random.normal)  
  
        self.dense2 = tf.keras.layers.Dense(1)  
  
    def call(self, x, training=True):  
        x = x[:, tf.newaxis]  
        x = self.dense1(x)  
        x = self.dense2(x)  
        return tf.squeeze(x, axis=1)
```

```
In [26]: model = Model(64)
```

```
In [27]: plt.plot(x.numpy(), y.numpy(), '.', label='data')
plt.plot(x, f(x), label='Ground truth')
plt.plot(x, model(x), label='Untrained predictions')
plt.title('Before training')
plt.legend();
```



Write a basic training loop:

```
In [28]: variables = model.variables


optimizer = tf.optimizers.SGD(learning_rate=0.01)

for step in range(1000):
    with tf.GradientTape() as tape:
        prediction = model(x)
        error = (y-prediction)**2
        mean_error = tf.reduce_mean(error)
    gradient = tape.gradient(mean_error, variables)
    optimizer.apply_gradients(zip(gradient, variables))

    if step % 100 == 0:
        print(f'Mean squared error: {mean_error.numpy():0.3f}')
```

```
Mean squared error: 10.085
Mean squared error: 1.083
Mean squared error: 1.077
Mean squared error: 1.074
Mean squared error: 1.071
Mean squared error: 1.068
Mean squared error: 1.066
Mean squared error: 1.065
Mean squared error: 1.064
Mean squared error: 1.062
```

```
In [29]: plt.plot(x.numpy(),y.numpy(), '.', label="data")
plt.plot(x, f(x), label='Ground truth')
plt.plot(x, model(x), label='Trained predictions')
plt.title('After training')
plt.legend();
```



That's working, but remember that implementations of common training utilities are available in the `tf.keras` module. So consider using those before writing your own. To start with, the `Model.compile` and `Model.fit` methods implement a training loop for you:

```
In [30]: new_model = Model(64)
```

```
In [31]: new_model.compile(
          loss=tf.keras.losses.MSE,
          optimizer=tf.optimizers.SGD(learning_rate=0.01))

history = new_model.fit(x, y,
                        epochs=100,
                        batch_size=32,
                        verbose=0)

model.save('./my_model')
```

INFO:tensorflow:Assets written to: ./my_model/assets

```
In [32]: plt.plot(history.history['loss'])  
plt.xlabel('Epoch')  
plt.ylim([0, max(plt.ylim())])  
plt.ylabel('Loss [Mean Squared Error]')  
plt.title('Keras training progress');
```

